

Gouki Minegishi

Ph.D. Candidate · The University of Tokyo

Tokyo, Japan
minegishi@weblab.t.u-tokyo.ac.jp
gouki510.github.io

RESEARCH INTERESTS

Mechanistic Interpretability

In-Context Learning

Large Language Models

Reasoning

Sparse Autoencoders

Neural Network Circuits

EDUCATION

Ph.D. in Computer Science

Apr 2025 – Mar 2028

The University of Tokyo · Advisor: Prof. Yutaka Matsuo

M.S. in Computer Science

Apr 2023 – Mar 2025

The University of Tokyo

B.S. in Engineering

Apr 2019 – Mar 2023

The University of Tokyo

SELECTED PUBLICATIONS

01 Emergent Analogical Reasoning in Transformers

Gouki Minegishi, Jingyuan Feng, Hiroki Furuta, Takeshi Kojima, Yusuke Iwasawa, and Yutaka Matsuo

International Conference on Machine Learning (ICML), 2026 **SPOTLIGHT (TOP 2.2%)**

02 Steering at the Source: Style Modulation Heads for Robust Persona Control

Yoshihiro Izawa, Gouki Minegishi, Koshi Eguchi, Sosuke Hosokawa, and Kenjiro Taura

International Conference on Machine Learning (ICML), 2026

03 Understanding Emergent Misalignment via Feature Superposition Geometry

Gouki Minegishi, Hiroki Furuta, Takeshi Kojima, Yusuke Iwasawa, and Yutaka Matsuo

Annual Meeting of the Association for Computational Linguistics (ACL), 2026

04 Mechanism of Task-oriented Information Removal in In-context Learning

Hakaze Cho, Haolin Yang, Gouki Minegishi, and Naoya Inoue

International Conference on Learning Representations (ICLR), 2026

05 RL Squeezes, SFT Expands: A Comparative Study of Reasoning LLMs

Kohsei Matsutani, Shota Takashiro, Gouki Minegishi, Takeshi Kojima, Yusuke Iwasawa, and Yutaka Matsuo

International Conference on Learning Representations (ICLR), 2026

06 LIT-RAGBench: Benchmarking Generator Capabilities of Large Language Models in Retrieval-Augmented Generation

Koki Itai, Shunichi Hasegawa, Yuta Yamamoto, Gouki Minegishi, and Masaki Otsuki

Joint International Conference on Computational Linguistics, Language Resources and Evaluation (LREC-COLING), 2026

07 In-Context Neurofeedback: Can LLMs Control Their Internal Representations through Privileged Access?

Koshiro Aoki, Ryota Takatsuki, Gouki Minegishi, Yusuke Haruki, and Daisuke Kawahara

Findings of the Association for Computational Linguistics: EMNLP, 2026

Trustworthy AI for Good (AI4GOOD) Workshop @ ICML, 2026 **SPOTLIGHT**

08 Zipping the Thought: When and How Compressed Reasoning Data Works in LLM Post-Training

Kohsei Matsutani, Gouki Minegishi, Takeshi Kojima, Yusuke Iwasawa, and Yutaka Matsuo

Findings of the Association for Computational Linguistics: EMNLP, 2026

09

Topology of Reasoning: Understanding Large Reasoning Models through Reasoning Graph Properties
Gouki Minegishi, Hiroki Furuta, Takeshi Kojima, Yusuke Iwasawa, and Yutaka Matsuo
Neural Information Processing Systems (NeurIPS), 2025

10

Beyond Induction Heads: In-Context Meta Learning Induces Multi-Phase Circuit Emergence
Gouki Minegishi, Hiroki Furuta, Shohei Taniguchi, Yusuke Iwasawa, and Yutaka Matsuo
International Conference on Machine Learning (ICML), 2025
ICLR Workshop on Building Trust in Language Models and Applications, 2025 ORAL

11

Rethinking Evaluation of Sparse Autoencoders through the Representation of Polysemous Words
Gouki Minegishi, Hiroki Furuta, Yusuke Iwasawa, and Yutaka Matsuo
International Conference on Learning Representations (ICLR), 2025

12

Interpreting Multi-Attribute Confounding through Numerical Attributes in Large Language Models
Hirohane Takagi*, Gouki Minegishi*, Shota Kizawa, Issey Sukeda, Hitomi Yanaka (* Equal Contribution)
International Joint Conference on NLP / Asia-Pacific Chapter of ACL (IJCNLP-AACL), 2025

13

Bridging Lottery Ticket and Grokking: Understanding Grokking from Inner Structure of Networks
Gouki Minegishi, Yusuke Iwasawa, and Yutaka Matsuo
Transactions on Machine Learning Research (TMLR)

14

Towards Empirical Interpretation of Internal Circuits and Properties in Grokked Transformers on Modular Polynomials
Hiroki Furuta, Gouki Minegishi, Yusuke Iwasawa, and Yutaka Matsuo
Transactions on Machine Learning Research (TMLR)
ICLR Workshop on Bridging the Gap Between Practice and Theory in Deep Learning, 2024 ORAL

15

ADOPT: Modified Adam Can Converge with Any β_2 with the Optimal Rate
Shohei Taniguchi, Keno Harada, Gouki Minegishi, et al.
Neural Information Processing Systems (NeurIPS), 2024

RESEARCH & WORK EXPERIENCE

Student Researcher	May 2026 – Oct 2026
Google DeepMind, Japan	
Research Scientist	Jun 2025 – Present
Third Intelligence, Inc.	
Research on developing new forms of intelligence beyond conventional AI, aiming to bridge human-like reasoning, perception, and action.	
Chief AI Engineer	Nov 2024 – May 2025
Matsuo Institute, Inc.	
Developed automated driving software, including 3D object detection with deep learning and integration of the World Model into automated driving.	
Research Intern	Aug 2024 – Oct 2024
Preferred Networks, Inc.	
Research on in-context learning in foundation models, with knowledge conflicts.	

ACADEMIC HONORS

Best Paper Award, JSAI 2026	2026
Spotlight Presentation (top 2.2%), ICML 2026	2026
Spotlight Presentation, ICML 2026 AI4GOOD Workshop	2026
Young Researcher Encouragement Award, NLP	2025
Outstanding Presentation Award, JSAI 2024	2024
Oral Presentation, ICLR 2025 Workshop (Building Trust in LMs)	2025
Oral Presentation, ICLR 2024 Workshop (Bridging Practice and Theory)	2024
Selected for "BOOST NAIS" Scholarship Program for fostering advanced AI talents	

■ ACADEMIC SERVICE

Reviewer:

- ICML — 2026 (Silver Reviewer), 2025
- NeurIPS — 2026, 2025
- ICLR — 2026
- AAAI — 2026
- TMLR

Organizer: Mechanistic Interpretability Organized Session at JSAI 2025